

BACHELOR INFORMATICA



UNIVERSITY OF AMSTERDAM

Creating an experimental environment for participating in TREC Fair Ranking Track

René Emmaneel

November 25, 2020

Supervisor(s): Fatemeh Sarvi

Signed:

Abstract

The TREC Fair Ranking Track is an annual academic search task with the stated goal of developing a benchmark for evaluating retrieval systems in terms of fairness. To participate we created a framework. [1]

Contents

1	Introduction	7
2	Theoretical background	9
2.1	Ethical aspect	9
2.2	Measuring fairness	9
2.2.1	Utility of a ranking	9
2.3	Learning to rank	9
2.3.1	Feature extraction	10
2.4	TREC Fair Ranking Track definition	11
2.4.1	Semantic Scholar Open Research Corpus	11
2.5	12
2.6	12
3	Creating the framework	13
4	Experiments	15
5	Conclusions	17
	Bibliography	19

Introduction

Ranking systems allow items to be seen by the users, where the exposure each item gets is significantly determined by the position in the ranking. Traditionally a ranking system would order items based on their relevance so that the system is of maximum utility to the user. However, modern information access systems often influence both the consumer and the producer of the content that they serve. Examples are hiring platforms where it is important that the system has a high utility for both the employer and the job-seeker. But also more traditional environments such as music, book or video recommendations can be considered two-sided because the indirect matching of users to content creators.

The ranking system allows exposure to the producers, which exposes risks to companies such as Google that hold a lot of power, and has already faced legal challenges such as the European Union antitrust violation fine [2]. While it is unlikely that a universal definition of fairness can be created, there are various attempts to measure unfairness and devise various fair ranking algorithms to ensure that the received attention of a given subject is approximately equal to its deserved attention [1].

To evaluate different fair ranking algorithms, the TREC Fair Ranking track was created. The stated goal was to develop a benchmark for evaluating retrieval systems in terms of fairness, as well as releasing a dataset for benchmarking fair ranking algorithms. The TREC 2020 Track was an academic search task, where a set of academic article abstract and queries were submitted to an academic search engine. The central goal is to provide fair exposure to different groups of authors, where the group definitions can be arbitrary. Various sets of data has been provided, as well as a detailed evaluation protocol. The data of the 2019 track includes the Semantic Scholar Open Corpus, which consists of data describing various papers. It also includes query data, and query sequences.

To participate in the TREC fair ranking track, an environment needs to be build for the entire process. This process includes handling of the input data, implementing the evaluation algorithm, making it modular for various uses, testing of the framework and implementing learning to rank machine learning algorithms. To design and implement the framework the following research question will be answered:

How can a modular framework be constructed to create an experimental environment for participating in the TREC fair ranking track.

TODO: Structure van project

Theoretical background

2.1 Ethical aspect

2.2 Measuring fairness

Fairness can be subjective and differ for various use cases. Before describing the fairness definition we have to describe the way exposure can influence the discoverability for a certain group.

2.2.1 Utility of a ranking

2.3 Learning to rank

Before looking at a way to create a model for learning to fairly rank, we will first propose a model for the standard learning to rank model on the TREC fair ranking track data. Learning to rank, when applied to document retrieval, is a task as follows. In training a set of queries and documents for each query is provided, as is the relevance judgement for each document. Using this data, a ranking function is created such that the model can predict a ranked list of documents for a query. The two major approaches for creating a ranking function is either a non-learning approach or a learning to rank approach. Learning approaches for information retrieval (IR) have been widely studied and shown to be useful for ranking articles. [3]

A learning to rank framework is systematically represented in Figure 2.1. As can be seen in the figure, a training set is needed to train the ranking model. The training data consist typically of n training queries, and m documents for each query. Each document is represented by a feature vector, which is a list of values calculated before training based on the document and query. Each query also has a relevance judgement y , which can for example be the order in which the documents are ranked, or can be a vector representing the relevance for each document based on the query. After training a model has been created, which can be used to rank a set of given documents based on the query.

There are three main learning to rank algorithms. The pointwise approach, which learns a function for directly predicting the relevance for each document for a query. The pairwise approach, which learns a function that takes in a pair of documents and a query and tries to predict the document that is most relevant for the given query. Lastly there is the listwise approach, which learns a function to directly rank a set of input documents based on a query.

In this thesis we have decided to create a pairwise ranking algorithm. This decision was made based on the training data available in the TREC fair ranking track. As described in [section], the training data consist of queries and a certain amount of documents per query, with a boolean relevance score for each document. By matching documents that are relevant to a certain query

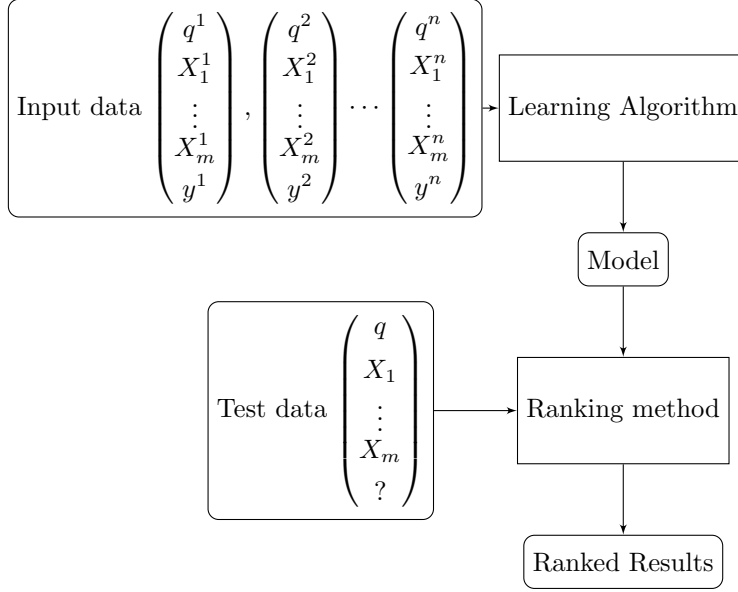


Figure 2.1: Learning to rank framework

ID	Feature Description
1	$\sum_{q_i \in q \cap d} TF(q_i, d)$ in title
2	$\sum_{q_i \in q} IDF(q_i)$ in title
3	$\sum_{q_i \in q \cap d} TF(q_i, d) \cdot IDF(Q_i)$ in title

Table 2.1: Features used in the rankSVM model

with documents that are irrelevant to the same query, we can easily create document pairs for the pairwise approach. By creating these document pairs, we can reduce the problem to a classification problem to decide which pair of the document is more relevant. We will first discuss the various features extracted for this model, after which we will discuss the learning algorithm itself.

2.3.1 Feature extraction

In learning to rank, each query-document pair is represented by a feature vector, and each feature gives an indication to the relevance of the document to the query. For example one feature could be the frequency of the query words in the abstract of the document. In this section we will discuss the various types of features, and which features we have used in constructing the feature vectors for the learning to rank framework.

Features can be categorized in 3 distinct categories. The query-document category means that the feature is dependant on both the feature and document, examples are term frequency and BM25. The other two categories are only dependant on the query or the document respectively.

The different features as used in the rankSVM model for the TREC fair ranking track are described in table 2.1. The explanation about these features are listed below.

- Features (A-B) are the term frequency (TF) features. The features will count the amount of occurrences of each term in the query in the respective field on the document.
- Features (C-D) are the inverse document frequency (IDF) features. The inverse document frequency function is a measure about how much information a given word provides. The function is defined as follows.

$$IDF(t) = \log \frac{N}{n(t)}$$

with N being the total number of documents, and $n(t)$ being the amount of documents with the term t

- Features (E-F) are the term frequency - inverse document frequency (TF-IDF) features. For every term of the query, the TF and IDF of the specific term is multiplied. This feature is correlated by a high frequency of a term, which has a low document frequency.
- Features (G-H) are the BM25 features. This is a TF-IDF-like retrieval function, first developed for TREC-3 [4]. The function used is defined as follows.

$$BM25(q, d) = \sum_{q_i \in q \cap d} IDF(q_i) \cdot \frac{TF(q_i, d) \cdot (k + 1)}{TF(q_i, d) + k \cdot (1 - b + b \cdot \frac{length(d)}{avlength()})}$$

where $length(d)$ is the amount of words in document d , and $avlength()$ is the average amount of words in all documents. k and b are free variables, which can be set to various values dependending on the dataset, but are set to $k=1.2$ and $b=0.75$, which experiments have shown to be reasonable values [5].

2.4 TREC Fair Ranking Track definition

2.4.1 Semantic Scholar Open Research Corpus

The Sementic Scholar Open Research Corpus is a large archive containing metadata of research papers [6]. It contains 47GB data in JSON format, containing information about papers in many different disciplines. A large dataset is needed for sufficient training sample size. The following data is available for most papers.

- **S2 Paper ID**
- **DOI**
- **Title**
- **Year**
- **Abstract**
- **Venue name**
- **Authors** (resolved to author IDs)
- **Inbound citations** (resolved to S2 paper IDs)
- **outbound citations** (resolved to S2 paper IDs)

Using this data, we can extract the corpus into smaller csv files containing the metadata of each paper and author. The TREC fair ranking track provided the following 3 csv files, however we also created functionality to extract the following information ourself given the Sementic Scholar Open Research Corpus. The paper metadata is as follows.

paper_metadata.csv:

- **S2 Paper ID**
- **Title**
- **Year**
- **Venue**
- **Number of inbound citations**

Using the same data, we can also extract metadata information for each author of the papers. During the extraction, we count for each author the amount of papers they worked on and the amount of citations each papers has. We use this information for calculating the i10 and H_index scores, as explained below.

author_metadata.csv:

- **Author ID**
- **Name**
- **Citation count:** sum of all citations of each published paper
- **Paper count**
- **i10:** amount of publications with atleast 10 citations
- **H_index:** The maximum value of h such that the given author has published h papers that have each been cited at least h times.
- **H_class:** The class is 'H' if the author has a H_index of atleast 10, and the class is 'L' if the H_index is less than 10

To combine the papers and authors, we need a third linker file, which combined the paper ID with each corresponding author ID, and the position in which they were attributed.

- **S2 Paper ID**
- **Author ID**
- **Position**

2.5

2.6

Creating the framework

Suspendisse vitae elit. Aliquam arcu neque, ornare in, ullamcorper quis, commodo eu, libero. Fusce sagittis erat at erat tristique mollis. Maecenas sapien libero, molestie et, lobortis in, sodales eget, dui. Morbi ultrices rutrum lorem. Nam elementum ullamcorper leo. Morbi dui. Aliquam sagittis. Nunc placerat. Pellentesque tristique sodales est. Maecenas imperdiet lacinia velit. Cras non urna. Morbi eros pede, suscipit ac, varius vel, egestas non, eros. Praesent malesuada, diam id pretium elementum, eros sem dictum tortor, vel consectetur odio sem sed wisi.

CHAPTER 4

Experiments

CHAPTER 5

Conclusions

Bibliography

- [1] Asia J. Biega, Krishna P. Gummadi, and Gerhard Weikum. Equity of attention: Amortizing individual fairness in rankings. *CoRR*, abs/1805.01788, 2018.
- [2] Mark Scott. Google fined record \$2.7 billion in e.u. antitrust ruling. *New York Times*, 2017.
- [3] Tao Qin, Tie-Yan Liu, Jun Xu, and Hang Li. Letor: A benchmark collection for research on learning to rank for information retrieval. August 2016.
- [4] Stephen E. Robertson; Steve Walker; Susan Jones; Micheline Hancock-Beaulieu Mike Gatford. Okapi at trec-3. 1994.
- [5] Hinrich Schütze Christopher D. Manning, Prabhakar Raghavan. An introduction to information retrieval. *Cambridge University Press*, page 270, 2009.
- [6]